

Lab4 实验报告

姓名：李雨朔 学号：20232131090 班级：计科一班 日期：2025.11.21

一 . 实验环境准备

```
git fetch  
git checkout fs  
make clean
```

二 . 实验 4.1 对 xv6 文件系统添加“大文件”的支持

- 在 kernel/fs.h 和 kernel/file.h 中修改宏定义和数据结构：

```
//kernel/fs.h  
#define NDIRECT 11           // 减少一个直接块用于二级间接块  
#define NINDIRECT (BSIZE / sizeof(uint)) // 每个间接块包含 256 个地址  
#define NTWOINDIRECT ((NINDIRECT)*(NINDIRECT)) // 二级间接块包含 65536 个地址  
#define MAXFILE (NDIRECT + NINDIRECT + NTWOINDIRECT)  
// 最大文件块数：11+256+65536=65803  
  
struct dinode {  
    short type;  
    short major;  
    short minor;  
    short nlink;  
    uint size;  
    uint addrs[NDIRECT+2]; // 11 个直接块 + 1 个一级间接块 + 1 个二级间接块  
};  
  
//kernel/file.h  
struct inode {  
    uint dev;  
    uint inum;  
    int ref;  
    struct sleeplock lock;  
    int valid;  
  
    short type;
```

```

    short major;
    short minor;
    short nlink;
    uint size;
    uint addrs[NDIRECT+2]; // 与磁盘 inode 结构保持一致
};


```

```

#define FSMAGIC 0x10203040
#define NDIRECT 11
#define NINDIRECT (BSIZE / sizeof(uint))
#define NTWOINDIRECT ((NINDIRECT)*(NINDIRECT))
#define MAXFILE (NDIRECT + NINDIRECT + NTWOINDIRECT)

// On-disk inode structure
struct dinode {
    short type;          // File type
    short major;         // Major device number (T_DEVICE only)
    short minor;         // Minor device number (T_DEVICE only)
    short nlink;         // Number of links to inode in file system
    uint size;           // Size of file (bytes)
    uint addrs[NDIRECT+2]; // Data block addresses
};


```

```

struct file {
    enum { FD_NONE, FD_PIPE, FD_INODE, FD_DEVICE } type;
    int ref; // reference count
    char readable;
    char writable;
    struct pipe *pipe; // FD_PIPE
    struct inode *ip; // FD_INODE and FD_DEVICE
    uint off; // FD_INODE
    short major; // FD_DEVICE
};

#define major(dev) ((dev) >> 16 & 0xFFFF)
#define minor(dev) ((dev) & 0xFFFF)
#define mkdev(m,n) ((uint)((m)<<16| (n)))

// in-memory copy of an inode
struct inode {
    uint dev;           // Device number
    uint inum;          // Inode number
    int ref;            // Reference count
    struct sleeplock lock; // protects everything below here
    int valid;          // inode has been read from disk?

    short type;         // copy of disk inode
    short major;
    short minor;
    short nlink;
    uint size;
    uint addrs[NDIRECT+2];
};

// map major device number to device functions.
struct devsw {
    int (*read)(int, uint64, int);
    int (*write)(int, uint64, int);
};

extern struct devsw devsw[];

#define CONSOLE 1

```

2. 在 kernel/fs.c 中修改 bmap 函数，添加二级间接块支持：

```

static uint
bmap(struct inode *ip, uint bn)
{
    uint addr, *a;
    struct buf *bp;

```

```

if(bn < NDIRECT)
{
    //获得直接块的地址, 若为 0 就分配
    if((addr = ip->addrs[bn]) == 0)
        ip->addrs[bn] = addr = balloc(ip->dev);
    return addr;
}

bn -= NDIRECT;
if(bn < NINDIRECT)
{
    //获得一级间接块的地址, 若为 0 就分配
    if((addr = ip->addrs[NDIRECT]) == 0)
        ip->addrs[NDIRECT] = addr = balloc(ip->dev);
    bp = bread(ip->dev, addr); //根据地址读一级间接块
    a = (uint*)bp->data; //一级间接块的数据指针
    if((addr = a[bn]) == 0)
    {
        //根据指针从相应偏移量处获得数据块地址, 若为 0 就分配
        a[bn] = addr = balloc(ip->dev);
        log_write(bp);
    }
    //获得地址后释放一级间接块并返回
    brelse(bp);
    return addr;
}

bn -= NINDIRECT;
struct buf *bp_2;
uint *a_2;
if(bn < NTWOINDIRECT)
{
    //获得二级间接块的地址, 若为 0 就分配
    if((addr = ip->addrs[NDIRECT+1]) == 0)
        ip->addrs[NDIRECT+1] = addr = balloc(ip->dev);
    bp_2 = bread(ip->dev, addr); //根据地址读二级间接块
    a_2 = (uint*)bp_2->data; //二级间接块的数据指针
    if((addr = a_2[bn/NINDIRECT]) == 0)
    {
        //根据指针从相应偏移量处获得一级间接块地址, 若为 0 就分配
        a_2[bn/NINDIRECT] = addr = balloc(ip->dev);
        log_write(bp_2);
    }
    //获得一级间接块地址后释放二级间接块
    brelse(bp_2);
    bp = bread(ip->dev, addr); //读一级间接块
}

```

```

a = (uint*)bp->data; //一级间接块的数据指针
if((addr = a[bn%NINDIRECT]) == 0)
{
    //根据指针从相应偏移量处获得数据块地址，若为0就分配
    a[bn%NINDIRECT] = addr = balloc(ip->dev);
    log_write(bp);
}
//获得数据块地址后释放一级间接块并返回
brelease(bp);
return addr;
}
panic("bmap: out of range");
}

```

```

if((addr = ip->addrs[bn]) == 0){
    //addr = balloc(ip->dev);
    ip->addrs[bn] = addr = balloc(ip->dev);
    //if(addr == 0)
    //    //return 0;
    //ip->addrs[bn] = addr;
}
return addr;
}
}
bn -= NDIRECT;
if(bn < NINDIRECT){
    // Load indirect block, allocating if necessary.
    if((addr = ip->addrs[NINDIRECT]) == 0){
        ip->addrs[NINDIRECT] = addr = balloc(ip->dev);
        //addr = balloc(ip->dev);
        //if(addr == 0)
        //    //return 0;
        //ip->addrs[NINDIRECT] = addr;
    }
    bp = bread(ip->dev, addr);
    a = (uint*)bp->data;
    if((addr = a[bn]) == 0){
        a[bn] = addr = balloc(ip->dev);
        //addr = balloc(ip->dev);
        //if(addr){
        //    //a[bn] = addr;
        log_write(bp);
    }
    brelease(bp);
}
bp = bread(ip->dev, addr);
a = (uint*)bp->data;
if((addr = a[bn%NINDIRECT]) == 0) {
    a[bn%NINDIRECT] = addr = balloc(ip->dev);
    log_write(bp);
}
brelease(bp);
return addr;
}

```

```

uint *a_2;
if(bn < NTWOINDIRECT) {
    //uint level1_Index = bn / NINDIRECT;
    //uint level2_Index = bn % NINDIRECT;
    if((addr = ip->addrs[NINDIRECT+1]) == 0){
        ip->addrs[NINDIRECT+1] = addr = balloc(ip->dev);
    }
    bp_2 = bread(ip->dev, addr);
    a_2 = (uint*)bp_2->data;
    if((addr = a_2[bn%NINDIRECT]) == 0) {
        a_2[bn/NINDIRECT] = addr = balloc(ip->dev);
        log_write(bp_2);
    }
    brelease(bp_2);
}
bp = bread(ip->dev, addr);
a = (uint*)bp->data;
if((addr = a[bn%NINDIRECT]) == 0) {
    a[bn%NINDIRECT] = addr = balloc(ip->dev);
    log_write(bp);
}
brelease(bp);
return addr;
}
//printf("bmap panic: bn=%d, NDIRECT=%d, NINDIRECT=%d, NDINDIRECT=%d\n", (int)(bn+NDIRECT+NINDIRECT), (int)NDIRECT, (int)
panic("bmap: out of range");
}

```

3. 在 kernel/fs.c 中修改 itrunc 函数，添加二级间接块的释放逻辑：

```
void
itrunc(struct inode *ip)
{
    int i, j;
    struct buf *bp,*bp_2;
    uint *a,*a_2;
    for(i = 0; i < NDIRECT; i++)
    {
        if(ip->addrs[i])
        {
            //遍历所有直接块地址，若不为 0 就释放然后设为 0
            bfree(ip->dev, ip->addrs[i]);
            ip->addrs[i] = 0;
        }
    }
    if(ip->addrs[NDIRECT])
    {
        //一级间接块不为 0
        bp = bread(ip->dev, ip->addrs[NDIRECT]); //读一级间接块
        a = (uint*)bp->data; //一级间接块的数据指针
        for(j = 0; j < NINDIRECT; j++)
        {
            //释放一级间接块下的所有数据块
            if(a[j])
                bfree(ip->dev, a[j]);
        }
        //释放一级间接块本身
        brelse(bp);
        bfree(ip->dev, ip->addrs[NDIRECT]);
        ip->addrs[NDIRECT] = 0;
    }
    if(ip->addrs[NDIRECT+1])
    {
        //二级间接块不为 0
        bp_2 = bread(ip->dev, ip->addrs[NDIRECT+1]); //读二级间接块
        a_2 = (uint*)bp_2->data; //二级间接块的数据指针
        for(i = 0; i < NINDIRECT; i++)
        {
            //遍历二级间接块下的所有一级间接块
            if(a_2[i])
            {
                //该一级间接块不为空
                bp = bread(ip->dev, a_2[i]);
            }
        }
    }
}
```

```

a = (uint*)bp->data;
for(j = 0; j < NINDIRECT; j++)
{
    if(a[j])
        bfree(ip->dev, a[j]);
}
//释放一级间接块本身
brelse(bp);
bfree(ip->dev, a_2[i]);
a_2[i] = 0;
}
}
//释放二级间接块本身
brelse(bp_2);
bfree(ip->dev, ip->addrs[NDIRECT+1]);
ip->addrs[NDIRECT+1] = 0;
}
ip->size = 0;
iupdate(ip);
}

```

```

void
ltruncate(struct inode *ip)
{
    int i, j;
    struct buf *bp, *bp2;
    uint *a, *a2;

    for(i = 0; i < NDIRECT; i++){
        if(ip->addrs[i]){
            bfree(ip->dev, ip->addrs[i]);
            ip->addrs[i] = 0;
        }
    }

    if(ip->addrs[NDIRECT]){
        bp = bread(ip->dev, ip->addrs[NDIRECT]);
        a = (uint*)bp->data;
        for(j = 0; j < NINDIRECT; j++){
            if(a[j])
                bfree(ip->dev, a[j]);
        }
        brelse(bp);
        bfree(ip->dev, ip->addrs[NDIRECT]);
        ip->addrs[NDIRECT] = 0;
    }

    if(ip->addrs[NDIRECT+1]) {
        bp2 = bread(ip->dev, ip->addrs[NDIRECT+1]);
        a2 = (uint*)bp2->data;
        for(i = 0; i<NINDIRECT; i++) {
            if(a2[i]) {
                bp = bread(ip->dev, a2[i]);
                a = (uint*)bp->data;
                for(j = 0; j<NINDIRECT; j++) {
                    if(a[j]) {
                        bfree(ip->dev, a[j]);
                    }
                }
                brelse(bp);
                bfree(ip->dev, a2[i]);
                a2[i] = 0;
            }
        }
        brelse(bp2);
        bfree(ip->dev, ip->addrs[NDIRECT+1]);
        ip->addrs[NDIRECT+1] = 0;
    }
}

```

```
    }
    ip->size = 0;
    iupdate(ip);
}
```

4. 编译和测试：

```
make clean  
make qemu  
  
$ bigfile
```

```
$ bigfile  
.  
.  
.  
wrote 65803 blocks  
reading bigfile  
.  
.  
bigfile done; ok
```

三 . 实验 4.2 符号链接

1. 添加系统调用号和声明：

```
// kernel/syscall.h  
// 添加系统调用号  
#define SYS_symlink 22  
  
//kernel/syscall.c  
// 在 extern 声明部分添加  
extern uint64 sys_symlink(void);  
// 在系统调用函数指针数组中添加  
[SYS_symlink] sys_symlink,  
  
//user/user.h  
// 在系统调用声明部分添加  
int symlink(const char *target, const char *path);  
  
//user/usys.pl  
# 在入口列表中添加  
entry("symlink");
```

```
// System call numbers
#define SYS_fork 1
#define SYS_exit 2
#define SYS_wait 3
#define SYS_pipe 4
#define SYS_read 5
#define SYS_kill 6
#define SYS_exec 7
#define SYS_stat 8
#define SYS_chdir 9
#define SYS_dup 10
#define SYS_getpid 11
#define SYS_sbrk 12
#define SYS_sleep 13
#define SYS_uptime 14
#define SYS_open 15
#define SYS_write 16
#define SYS_mknod 17
#define SYS_unlink 18
#define SYS_link 19
#define SYS_mkdir 20
#define SYS_close 21
#define SYS_symlink 22
```

```
extern uint64 sys_fork(void);
extern uint64 sys_exit(void);
extern uint64 sys_wait(void);
extern uint64 sys_pipe(void);
extern uint64 sys_read(void);
extern uint64 sys_kill(void);
extern uint64 sys_exec(void);
extern uint64 sys_fstat(void);
extern uint64 sys_chdir(void);
extern uint64 sys_dup(void);
extern uint64 sys_getpid(void);
extern uint64 sys_sbrk(void);
extern uint64 sys_sleep(void);
extern uint64 sys_uptime(void);
extern uint64 sys_open(void);
extern uint64 sys_write(void);
extern uint64 sys_mknod(void);
extern uint64 sys_unlink(void);
extern uint64 sys_link(void);
extern uint64 sys_mkdir(void);
extern uint64 sys_close(void);
extern uint64 sys_symlink(void);
```

```
static uint64 (*syscalls[])() = {  
    [SYS_fork] sys_fork,  
    [SYS_exit] sys_exit,  
    [SYS_wait] sys_wait,  
    [SYS_pipe] sys_pipe,  
    [SYS_read] sys_read,  
    [SYS_kill] sys_kill,  
    [SYS_exec] sys_exec,  
    [SYS_fstat] sys_fstat,  
    [SYS_chdir] sys_chdir,  
    [SYS_dup] sys_dup,  
    [SYS_getpid] sys_getpid,  
    [SYS_sbrk] sys_sbrk,  
    [SYS_sleep] sys_sleep,  
    [SYS_uptime] sys_uptime,  
    [SYS_open] sys_open,  
    [SYS_write] sys_write,  
    [SYS_mknod] sys_mknod,  
    [SYS_unlink] sys_unlink,  
    [SYS_link] sys_link,  
    [SYS_mkdir] sys_mkdir,  
    [SYS_close] sys_close,  
    [SYS_symlink] sys_symlink,  
};
```

```
// system calls
int fork(void);
int exit(int) __attribute__((noreturn));
int wait(int*);
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(const char*, char**);
int open(const char*, int);
int mknod(const char*, short, short);
int unlink(const char*);
int fstat(int fd, struct stat*);
int lstat(const char*, const char*);
int mkdir(const char*);
int chdir(const char*);
int dup(int);
int getpid(void);
char* sbrk(int);
int sleep(int);
int uptime(void);
int symlink(const char *target, const c
```

```
entry("fork");
entry("exit");
entry("wait");
entry("pipe");
entry("read");
entry("write");
entry("close");
entry("kill");
entry("exec");
entry("open");
entry("mknode");
entry("unlink");
entry("fstat");
entry("link");
entry("mkdir");
entry("chdir");
entry("dup");
entry("getpid");
entry("sbrk");
entry("sleep");
entry("utime");
entry("symlink");
```

2. 定义符号链接类型和标志位:

```
//kernel/stat.h
#define T_SYMLINK 4 // 符号链接文件类型

//kernel/fcntl.h
#define O_NOFOLLOW 0x800 // 不跟随符号链接
```

```
#define T_DIR 1 // Directory
#define T_FILE 2 // File
#define T_DEVICE 3 // Device
#define T_SYMLINK 4

struct stat {
    int dev; // File system's disk device
    uint ino; // Inode number
    short type; // Type of file
    short nlink; // Number of links to file
    uint64 size; // Size of file in bytes
};
```

```
#define O_RDONLY 0x000
#define O_WRONLY 0x001
#define O_RDWR 0x002
#define O_CREAT 0x200
#define O_TRUNC 0x400
#define O_NOFOLLOW 0x800
```

3. 在 kernel/sysfile.c 文件中添加 sys_symlink 函数:

```
uint64
sys_symlink(void)
{
    char target[MAXPATH];
    char path[MAXPATH];
    struct inode *ip;
    // 从用户空间获取参数
    if(argstr(0, target, MAXPATH) < 0 || argstr(1, path, MAXPATH) < 0){
        return -1;
    }
```

```

// 开始文件系统操作
begin_op();
// 创建符号链接文件
ip = create(path, T_SYMLINK, 0, 0);
if(ip == 0)
{
    end_op();
    return -1;
}
// 将目标路径写入符号链接文件的数据块中
if(writei(ip, 0, (uint64)target, 0, MAXPATH) < MAXPATH)
{
    // 写入失败，需要清理
    iunlockput(ip);
    end_op();
    return -1;
}
iunlockput(ip);
end_op();
return 0;
}

```

```

uint64
sys_symlink(void)
{
    char target[MAXPATH];
    char path[MAXPATH];
    struct inode *ip;
    if(argstr(0, target, MAXPATH) < 0 || argstr(1, path, MAXPATH) < 0) {
        return -1;
    }
    begin_op();
    ip = create(path, T_SYMLINK, 0, 0);
    if(ip == 0) {
        end_op();
        return -1;
    }
    if(writei(ip, 0, (uint64)target, 0, MAXPATH) < MAXPATH) {
        iunlockput(ip);
        end_op();
        return -1;
    }
    iunlockput(ip);
    end_op();
    return 0;
}

```

4. 在 kernel/sysfile.c 的 sys_open 函数中添加符号链接处理逻辑：

```

// 符号链接处理：如果文件是符号链接且没有设置 O_NOFOLLOW 标志
if(ip->type == T_SYMLINK && !(omode & O_NOFOLLOW))
{
    // 这是个符号链接文件并且确定要 follow
    int depth = 10; // 最大递归深度
}

```

```
char target[MAXPATH];
for(int i = 0; i < depth; i++)
{
    // 从符号链接文件的数据块中读取出路径 target
    if(readi(ip, 0, (uint64)target, 0, MAXPATH) != MAXPATH)
    {
        iunlockput(ip);
        end_op();
        return -1;
    }
    // 原本的符号链接文件已经不需要了，释放它的锁
    iunlockput(ip);
    // 得到了 target，解析路径得到目标文件的 inode
    if((ip = namei(target)) == 0)
    {
        end_op();
        return -1;
    }
    ilock(ip);
    // 如果目标文件不再是符号链接，停止循环
    if(ip->type != T_SYMLINK)
        break;
    // 检查是否超过最大递归深度
    if(i == depth - 1)
    {
        iunlockput(ip);
        end_op();
        return -1;
    }
}
```

```

uint64
sys_open(void)
{
    char path[MAXPATH];
    int fd, omode;
    struct file *f;
    struct inode *ip;
    int n;

    argint(1, &omode);
    if((n = argstr(0, path, MAXPATH)) < 0)
        return -1;

    begin_op();

    if(omode & O_CREATE){
        ip = create(path, T_FILE, 0, 0);
        if(ip == 0){
            end_op();
            return -1;
        }
    } else {
        if((ip = namei(path)) == 0){
            end_op();
            return -1;
        }
        ilock(ip);

        if(ip->type == T_SYMLINK && !(omode & O_NOFOLLOW)) {
            int depth = 10;
            char target[MAXPATH];
            for(int i=0; i<depth; i++) {
                if(readi(ip, 0, (uint64)target, 0, MAXPATH) != MAXPATH) {
                    unlockput(ip);
                    end_op();
                    return -1;
                }
                unlockput(ip);
                if((ip = namei(target)) == 0) {
                    end_op();
                    return -1;
                }
                ilock(ip);
                if(ip->type!=T_SYMLINK) break;
            }
            if(i == depth-1) {
                unlockput(ip);
                end_op();
                return -1;
            }
        }
        if(ip->type == T_DIR && omode != O_RDONLY){
            unlockput(ip);
            end_op();
            return -1;
        }
    }
}

```

5. 修改 Makefile, 在 UPROGS 部分添加:

```
$U/_symlinktest\
```

```

UPROGS=\
$U/_cat\
$U/_echo\
$U/_forktest\
$U/_grep\
$U/_init\
$U/_kill\
$U/_ln\
$U/_ls\
$U/_mkdir\
$U/_rm\
$U/_sh\
$U/_stressfs\
$U/_usertests\
$U/_grind\
$U/_wc\
$U/_zombie\
$U/_symlinktest\_
```

6. 编译和测试：

```
make clean
```

```
make qemu
```

```
$ symlinktest
```

```
xv6 kernel is booting
init: starting sh
$ symlinktest
Start: test symlinks
test symlinks: ok
Start: test concurrent symlinks
test concurrent symlinks: ok
$
```

四 . 实验主观心得

本次实验中，在和 AI 智斗了很久后，终于走出 AI，然后发现 AI 世界外的代码清晰很多（这次的实验代码都来自于网上的大佬，自己的代码基础实在是太差了，基本很难独立写出代码）。AI 给我提供的很多解决方案在理论上可行，但实际上却无法出结果。然后去网上一查，再结合 AI 解释，发现好几个问题其实用简单的思路就能解决（大佬写得代码还清晰易懂，AI 的代码就没这种效果）。

仍然感谢 AI 的帮助。